


Overall Structure



The flowchart starts with a single point, branches out into multiple parallel paths based on a central decision step, and then converges back to a single endpoint.

Step-by-Step Flow

1. **Start:** The process begins at an oval node labeled "**Start.**"
2. **Intent Classification:** An arrow points from "Start" to a central rectangular node labeled "**Intent Classification.**" This node acts as the main sorting hub. It analyzes the user's request to understand what they want.
3. **Parallel Processes:** From "Intent Classification," several arrows point to different rectangular nodes, indicating the possible paths the application can take. These paths are:
 - **User Profiling:** This path is for gathering or updating information about the user.

- **Stock Data and Advice:** This path handles requests for stock market information and recommendations.
 - **Mock Expense Tracking:** This path is for simulating and tracking user expenses.
 - **Mock Budget Summary:** This path provides a simulated summary of the user's budget.
 - **Financial Advice:** This path is specifically for providing financial advice. There's a handwritten note next to this arrow that says "# for high risk handling," suggesting this path manages sensitive or high-stakes queries.
 - **HITL (Human in the Loop):** This path directs the query to a human agent for assistance.
 - **Fallback:** An arrow labeled "**Irrelevant**" points from "Intent Classification" to this node. This path is taken when the user's request is not related to finance.
4. **End:** All the paths listed above—from User Profiling to Fallback—have arrows pointing to a final oval node labeled "**End**." This signifies the completion of the user's request, regardless of which path was taken.